

Torchlight: Diffusion-based Network Trace Generation from the DARPA Searchlight Dataset

RAY ZHAO, University of Southern California,
ALEFIYA HUSSAIN, USC/ISI,

At present, there is a severe lack of both comprehensive and realistic labeled datasets for machine learning applications in the networking domain. Predominantly, prior generative work has focused on lower-dimensional representations that rely on aggregating flow characteristics and lack the fine-grain of raw network traces. This results in suboptimal performance in machine learning contexts and limited applications outside of those contexts. This has induced a push for new generative techniques to provide synthetic data for usage both on its own and layered in with real data as augmentation. In this paper, we present Torchlight, a diffusion-based generation framework built atop and extending techniques first introduced in *NetDiffusion* [?] using DARPA Searchlight [?] data to generate synthetic network traces for video streaming applications. We demonstrate the efficacy of Torchlight in generating synthetic network traces that reasonably resemble real-world data and perform notably well in classification tasks.

1 INTRODUCTION

There's a large need and general scarcity of labeled network datasets, as high-quality data is often expensive and fraught with privacy concerns, and the datasets that do exist are often left rarely update over time. Synthetic data generation techniques aim to solve this problem, but prior GAN based methodologies have been limited in their ability to generate raw, high-dimensional traces and often rely on aggregating flow characteristics, which results in data that lacks statistical fidelity to real data and performs much more poorly in machine learning tasks. The highly specific formatting of these types of condensed traces also renders them generally incompatible with traditional tools like Wireshark and tcpdump.

In recent years, the advent of diffusion models have solved many of the problems that plague older GAN frameworks, being able to capture much more complex patterns and relationships. They also usually offer much more stability in the training process, in which GANs are notably finicky. Additionally, their now near-hegemony in the image generation space has enabled a wide base of development support in the diffusion community and has also made available many "plug and play" means of constraining generation to appropriate enough scopes for specific generation tasks like the one at hand.

As it pertains to Torchlight, we build on the NetDiffusion framework pioneered by Jiang et al. [?] to generate synthetic network traces for video streaming applications, using the DARPA Searchlight dataset [?] as our base. Our contributions are as follows:

- (1) **Successful application of NetDiffusion techniques on the Searchlight Dataset:** Using the pre-processing and post-processing code provided by Jiang et al., we were able to successfully generate synthetic traces that performed capably on its own as training data for classification tasks as well as when used in conjunction with real data as augmentation.
- (2) **Streamlining the generation process:** While the original NetDiffusion codebase used externally-built WebUIs for LoRA fine-tuning as well as the Controlnet, these two major components were integrated into the codebase such that the entire process can be run from a single script.

2 MOTIVATION & RELATED WORK

Publicly available network data is crucial in this field for enabling continued research and development of new networking technologies and techniques. While useful for more traditional networking analysis tasks, machine learning applications have been driving much of the development in use-cases like anomaly detection and traffic classification and in particular have largely relied on these datasets.

2.1 Cost and Scarcity

Notable datasets like DARPA Searchlight and CAIDA have been essential in the development of new networking technologies, but expensive and complex to develop and then maintain. Entities with the means to develop these datasets often have little incentive to release them to the public, as they are often highly proprietary and run into ethics concerns. Even when they are released, they are often left to stagnate and lose relevance over time, as the networking landscape is highly dynamic.

2.2 Data Augmentation

In the absence of high-quality labeled data, data augmentation has been a popular technique, especially in computer vision, to extend the utility of existing and the models with which they are used. Synthetic data generation has been successful across domains, with GANs being dominant in the image generation space and autoencoders being popular in creating synthetic text data.

The most notable application of synthetic data generation in networking has been in the realm of flow-based data, where the data is aggregated into flows and then used to generate synthetic data. This has been successful in some applications, most notably NetShare [?], which produces NetFlow-like aggregates on network traffic. The limited dimensionality of these aggregates, however, means much coarser-grained pattern capture and limited utility outside of pure model training.

The original authors of NetDiffusion [?] have shown that NetFlow data performs poorly in classification tasks in comparison to raw traffic data, and as such have been motivated to develop a new framework that can generate raw traces.

3 METHODS

Here we outline the framework used in Torchlight to generate synthetic network traces, which is built atop NetDiffusion [?] and updates and streamlines it. Our approach is segmented into five primary moving parts:

- (1) Conversion of raw network traces in .pcap format into images
- (2) Fine-tuning a pretrained diffusion model with LoRA on the labelled images
- (3) ControlNet constrained generation of synthetic traces
- (4) Color correction and network compliance post-processing
- (5) Conversion of generated and processed images back into .pcap format

3.1 Trace to Image Conversion

This work is anchored on the core concept of encapsulating raw network data into a dense image format. Pixels in images inherently form high-dimensional patterns and structures with long-range, highly non-linear dependencies. While tabular formats typically engender points as independent entities, images much more adeptly preserve temporal and spatial relationships between network flows that traditional tabular formats do not. Development in the image space has always been rapid and deep-learning and subsequent generative techniques have been refined over time to capture these complex relationships. As importantly, this format allows us to exploit the maturity

and computational throughput of image-based deep learning. Diffusion models are the newest and most powerful of these techniques, and have been shown to be of great utility in the image generation space. We will delve into more detail on this in the next section.

Our image representation is expressed with the nPrint [?] format, which standardizes the irregular and high-dimensional network patterns of raw traces into a consistent bit-mapped format, where the presence or absence of a bit in the packet header is denoted 1 or 0 respectively, and a missing header bit is represented as -1. The payload content is not encoded since it's often encrypted, but size of payloads can be inferred from other encoded header fields. With this, a sequence of pcaps is converted into a matrix which can be formatted as an image. A set bit (1), an unset bit (0), and a vacant bit (-1) are represented by green, red, and blue respectively. Due to the limitations in generative models' capability to handle very high dimensional data and standard image processing practice, we group packets in groups of 1024, resulting in a 1088×1024 image with each row of pixels representing a packet in the network traffic flow as shown in Figure 2.

3.2 Model Fine-Tuning

3.2.1 Why diffusion models? Diffusion models fundamentally operate on the principle of iteratively updating a set of latent variables to approximate a true data distribution, taking a randomly initialized amount of noise and progressively refining it to generate an image. A differential equation is used to control the transformation of an initial noise vector z into a data point x over a series of time steps. This process is modeled by estimating the gradient of the log-likelihood of the data with respect to the noise vector, and then using this gradient to update the noise vector. Diffusion models typically operate in the context of "text-to-image" synthesis, where a text prompt guides the generation of an image by conditioning the noise vector on the text. This conditioning steers the data distribution from a general uniform prior distribution to a more specific one fitted to the text prompt.

This process lends a number of advantages to diffusion, as opposed to other generative techniques. Diffusion models have been shown to be adept at capturing highly complex and intricate relationships and generating high-resolution images. The text conditioning also allows for a high degree of control over the generated images, as the text can be used to guide the generation process in a way that is not possible with other generative models. Lastly, diffusion models are highly stable and easy to train, as they do not require the adversarial training that GANs do, and are much less prone to mode collapse. Their stable gradient dynamics and ability to be seeded in training ensures not just consistent but reproducible training results.

3.2.2 Fine-tuning with LoRA. Training a diffusion model from scratch is highly computationally intensive and expensive, especially considering that existing models are often trained on billions of images. While good for general applications, base models like Stable Diffusion are too broad for our purposes. While a text prompt like "pixelated network data: dash, 1080" is what we end up using in our end-stage model, the base model would produce something like a pixelated image of a dash, which is not useful for our purposes. Low-Rank Adaptation (LoRA) is a technique that allows us to fine-tune a pre-trained model on a smaller dataset, which is much more computationally feasible. LoRA works by attaching a low-rank matrix to the model's parameters and then training just the low-rank matrix on a new dataset. This allows us to adapt the stylistic and color characteristics of the model to our specific dataset without having to retrain the entire model.

In the fine-tuning process, we attach to every image a caption that corresponds to the various characteristics of the image. For example, the image for a network flow using the Wireguard VPN streaming 1080p video using DASH over HTTP3 would be captioned "pixelated network data: wireguardptp, dash, 1080 http3". Hundreds of such image-caption pairs are then used to fine-tune

the model over the course of thousands of epochs to adapt the model to the specific color and stylistic characteristics of the images used to fine-tune it. The choice of caption is crucial and achieves two main objectives: it gives the model a specific vocabulary that reduces ambiguity and minimizes interference from the original word embeddings of the base model.

3.3 Constrained Generation

Even though fine-tuning can result in very high-quality output, it can be very abstract and deviate greatly from the stricter bounds of network traffic data. To constrain the output, we can employ a ControlNet, which is a separate neural network that sits atop the diffusion model and modifies the output to be rigidly within the bounds defined by the ControlNet. Generally, a ControlNet is trained on a specific type of spatial constraint, like poses, depth mapping, masks, or in our case, canny edge detection. In the case of canny edge detection, the ControlNet takes as input a text prompt and the canny edges extracted from a sample image to use as "guardrails" to guide the diffusion model in generating the final image. ControlNet enables us to generate images that are not just high-quality, but also highly compliant with the structure of network traffic.

insert controlnet figure here

3.4 Post-Processing

While the constrained diffusion model generates high-fidelity images that greatly *resemble* network traffic, use outside of ML tasks require that the flows these images represent comply with network standards and the multitudes of inter and intra packet dependences that enable use in traditional networking applications. While ControlNet can do a lot to ensure resemblance, it cannot guarantee compliance. To this end, we must apply a series of post-processing steps to the generated images to ensure that they are compliant with network standards. Broadly, these rules can be put into two categories:

- (1) **Inter-Packet:** Rules that govern the sequencing between header fields between packets in a flow, like the SYN / SYN-ACK handshake that initiates a TCP connection, the incorrect alignment of which can cause the connection to fail.
- (2) **Intra-Packet:** Rules that organize the structure of individual packets, like the ordering of header fields within a packet and checksums that ensure the integrity of the packet.

Most of the work of creating a systematic approach that comprehensively identifies and corrects these errors has been done by original NetDiffusion authors, and we have largely built on their work in this area. The post-processing steps are fully automated and accomplished by traversing two intra-packet and inter-packet dependency trees built with domain knowledge. A notable challenge the authors faced was time-stamp assignment for packets, which was resolved by sampling original timestamps from real data to produce similar timestamp distributions in the synthetic data. This post-processing ensures that the generated traffic can be smoothly converted into raw network formats and utilized in non-ML contexts.

3.5 Image to Trace Conversion

The final step in the process is to convert the generated images back into network traces and is relatively straightforward. This is done by reversing the process used to convert the original traces into images. The generated images are first converted into nPrint format, and then the nPrint format is converted back into a sequence of pcaps. The post-processing steps ensure that the generated pcaps are compliant with network standards and can be used in traditional networking applications.

4 EVALUATION

The assessment of this framework is done primarily in gauging its performance in machine learning classification tasks. While the original NetDiffusion paper goes into great detail about the performance of the generated in outside applications, we focus on the classification task as it is the most common use-case for synthetic data in networking.

4.1 Dataset Overview & Generation

The data used is the DARPA Searchlight dataset [?], which is a large collection of network traces from a variety of sources and applications. Searchlight was collected over a period of 6 months and contains a wide variety of traffic types, including video streaming, web browsing, and file transfers all coordinated on a complex testbed that includes a variety of network devices and topologies, making it highly representative of real-world network traffic. We focus on video streaming traffic in our work.

4.2 Classification Results

ACKNOWLEDGMENTS

To Alefiya, for the continual guidance and enthusiasm all throughout the project

REFERENCES